

Transformation and differential abundance analysis of microbiome data incorporating phylogeny

Emily Palmer


Oregon State University

palmerem@oregonstate.edu

November 22, 2021

Genome analysis

Transformation and differential abundance analysis of microbiome data incorporating phylogeny

Chao Zhou^{1,2}, Hongyu Zhao^{2,3,*} and Tao Wang ^{1,2,4,*}

¹Department of Bioinformatics and Biostatistics, Shanghai Jiao Tong University, 200240 Shanghai, China, ²SJTU-Yale Joint Center for Biostatistics and Data Science, Shanghai Jiao Tong University, 200240 Shanghai, China, ³Department of Biostatistics, Yale University, New Haven, CT 06511, USA and ⁴MoE Key Lab of Artificial Intelligence, Shanghai Jiao Tong University, 200240 Shanghai, China

Outline: Overview of paper

- Extend Dirichlet-tree multinomial to zero-inflated Dirichlet tree multinomial to model multivariate microbial counts.
- Use a model-based posterior mean transformation to convert raw counts into non-zero relative abundances that sum to 1.
- Introduce phylogenetically informed DA procedure on transformed data - adaptive analysis of composition of microbiomes for DA testing (adaANCOM)- log-ratios adaptively on the tree for each taxon.

Introduction

- Differential abundance tests need to account for the high dimensional, sparse, compositional, negatively and positively correlated, phylogenetically structured microbiome data.
- Scaling methods to correct for compositional bias (TMM, DESeq, CSS) - assume that most taxa are not differentially abundant, when the count matrix is sparse it can over or underestimate diversity, distort correlations

Review of Dirichlet-Multinomial distribution

- The Dirichlet-multinomial (DM) distribution is commonly used for OTU counts and is a compound distribution comprised of Dirichlet and Multinomial distributions
- Let $\mathbf{y} = (y_1, \dots, y_K)^T$ be the count vector for a sample with K OTUs.
- $y^+ = \sum_{k=1}^K y_k$, $\mathbf{p} = (p_1, \dots, p_K)^T$, $\sum_{k=1}^K p_k = 1$, $\alpha^+ = \sum_{k=1}^K \alpha_k$
- Multinomial pdf:
- Dirichlet pdf:

$$f_M(\mathbf{y}; \mathbf{p}) = \frac{\Gamma(y^+ + 1)}{\prod_{k=1}^K \Gamma(y_k + 1)} \prod_{k=1}^K p_k^{y_k},$$

$$f_D(\mathbf{p}; \boldsymbol{\alpha}) = \frac{\Gamma(\alpha^+)}{\prod_{k=1}^K \Gamma(\alpha_k)} \prod_{k=1}^K p_k^{\alpha_k},$$

- Combining both we get the Dirichlet-Multinomial distribution:

$$\begin{aligned} f_{DM}(\mathbf{y}; \boldsymbol{\alpha}) &= \int f_M(\mathbf{y}; \mathbf{p}) f_D(\mathbf{p}; \boldsymbol{\alpha}) d\mathbf{p} \\ &= \frac{\Gamma(y^+ + 1) \Gamma(\alpha^+)}{\Gamma(y^+ + \alpha^+)} \prod_{k=1}^K \frac{\Gamma(y_k + \alpha_k)}{\Gamma(y_k + 1) \Gamma(\alpha_k)}. \end{aligned}$$

Relationship between Gamma, Beta, & Dirichlet distributions

- Let Z_1, \dots, Z_K independent $Z_k \sim \text{Gamma}(\alpha_k, \lambda)$,
- $X_k = \frac{Z_k}{\sum_{j=1}^K Z_j}$
- $W_k = \frac{Z_k}{\sum_{j=k}^K Z_j}$
- Joint distribution of $\mathbf{X} = (X_1, \dots, X_K)^T$ is Dirichlet distributed with parameter $\boldsymbol{\alpha}$
- $W_k \sim \text{Beta}(\alpha_k, \alpha_k^+)$ independent
- $X_k = W_k \prod_{j=1}^{k-1} (1 - W_j)$
- We can rewrite the Dirichlet distribution as

$$f_{DM}(\mathbf{y}; \boldsymbol{\alpha}) = \int f_M(\mathbf{y}; h(\mathbf{w})) \prod_{k=1}^{K-1} f_B(w_k; \alpha_k, \alpha_k^+) dw.$$

- $h(\cdot)$ is the transformation from \mathbf{W} to \mathbf{X} .

Zero Inflated Dirichlet Multinomial

- We can introduce zero inflation by defining the Zero Inflated Dirichlet Multinomial (ZIDM) as:

$$f_{ZIDM}(\mathbf{y}; \boldsymbol{\pi}, \boldsymbol{\alpha}) = \int f_M(\mathbf{y}; h(\mathbf{w})) \prod_{k=1}^{K-1} f_{ZIB}(w_k; \pi_k, \alpha_k, \alpha_k^+) d\mathbf{w}$$

- where

$$f_{ZIB}(w_k; \pi_k, \alpha_k, \alpha_k^+) = \pi_k \delta(0) + (1 - \pi_k) f_B(w_k; \alpha_k, \alpha_k^+).$$

is a zero-inflated binomial distribution

- π_k is the probability of zero-inflation in the k th component $\delta(\cdot)$ is the Dirac delta function

Introducing the Dirichlet Tree Multinomial

- Suppose the relationships between OTUs are encoded in a tree \mathcal{T} which is composed of internal nodes \mathcal{V} and leaf nodes \mathcal{L}
- Leaf nodes are OTUs, internal nodes are the node to branch on the phylogenetic tree
- For each $v \in \mathcal{V}$, denote C_v as the set of child nodes of v , \mathbf{y}_v the vector of counts corresponding to C_v , and $y_v^+ = \sum_{u \in C_v} y_u$.

Dirichlet Tree Multinomial & Zero-Inflated Dirichlet Tree Multinomial

- Assume that \mathbf{y}_v conditional on y_v^+ are independent across internal nodes, the Dirichlet Tree Multinomial (DTM) distribution is the product of DM distributions that factorize over the tree

$$f_{DTM}(\mathbf{y}; \boldsymbol{\alpha}_v, v \in \mathcal{V}) = \prod_{v \in \mathcal{V}} f_{DM}(\mathbf{y}_v; y_v^+, \boldsymbol{\alpha}_v)$$

- To extend the DTM to the ZIDTM, we replace the Dirichlet multinomial with Zero-Inflated Dirichlet Multinomial

$$f_{ZIDTM}(\mathbf{y}; \boldsymbol{\pi}_v, \boldsymbol{\alpha}_v, v \in \mathcal{V}) = \prod_{v \in \mathcal{V}} f_{ZIDM}(\mathbf{y}_v; y_v^+, \boldsymbol{\pi}_v, \boldsymbol{\alpha}_v)$$

Example




Fig. 1. A binary tree with $K=10$ leaves. Here $\mathcal{L} = \{1, 2, \dots, 10\}$, $\mathcal{V} = \{11, 12, \dots, 19\}$. For illustration, $\mathcal{C}_{17} = \{18, 19\}$, $\mathbf{y}_{17} = (y_{18}, y_{19})^T$ and $y_{17}^+ = y_{18} + y_{19}$. Given y_{17}^+ , \mathbf{y}_{17} has a DM or ZIDM distribution. The factorization over internal nodes means that these conditional distributions are independent

Maximum likelihood estimation for ZIDTM

- Use EM-algorithm to estimate unknown parameters:
 $\theta = \{\alpha_v, \pi_v, v \in \mathcal{V}\}$
- Assume $C_v = \{1, \dots, K_v\}$, where $K_v = |C_v|$
- Consider n observations $\mathbf{y}^1, \dots, \mathbf{y}^n$, then the log-likelihood function (without constant terms) is:

$$l(\theta) = \sum_{i=1}^n \sum_{v \in \mathcal{V}} \sum_{k=1}^{K_v-1} \{l_1(\delta_{vk}^i, \pi_{vk}) + l_2(\delta_{vk}^i, w_{vk}^i, \alpha_{vk}, \alpha_{vk}^+)\}$$

with

$$l_1(\delta_{vk}^i, \pi_{vk}) = \delta_{vk}^i \log \pi_{vk} + (1 - \delta_{vk}^i) \log(1 - \pi_{vk})$$

$$l_2(\delta_{vk}^i, w_{vk}^i, \alpha_{vk}, \alpha_{vk}^+) = (1 - \delta_{vk}^i) \{-\log \mathcal{B}(\alpha_{vk}, \alpha_{vk}^+) \\ + (\alpha_{vk} - 1) \log w_{vk}^i + (\alpha_{vk}^+ - 1) \log(1 - w_{vk}^i)\}$$

- δ_{vk}^i is the indicator of zero-inflation $\mathcal{B}()$ is the beta function

- Compute expectation of $l(\boldsymbol{\theta})$ with respect to posterior distribution of $(\delta_{vk}^i, w_{vk}^i) | \mathbf{y}_v^i$, indexed by current value of $\boldsymbol{\theta}$, which gives the Q -function:

$$Q(\boldsymbol{\theta}) = \sum_{i=1}^n \sum_{v \in \mathcal{V}} \sum_{k=1}^{K_v-1} E\{l_1(\delta_{vk}^i, \pi_{vk}) + l_2(\delta_{vk}^i, w_{vk}^i, \alpha_{vk}, \alpha_{vk}^+)\}$$

E-step: Definitions

- $\delta_{vk}^{i*} = E(\delta_{vk}^i | \mathbf{y}_v^i)$
- $R_{vk}^{i*} = E(\log w_{vk}^i | \mathbf{y}_v^i, \delta_{vk}^i = 0)$
- $S_{vk}^{i*} = E\{\log(1 - w_{vk}^i) | \mathbf{y}_v^i, \delta_{vk}^i = 0\}$
- Then,

$$\delta_{vk}^{i*} = \begin{cases} 0, & y_{vk}^i > 0, \\ \frac{\pi_{vk}}{\pi_{vk} + (1 - \pi_{vk}) \frac{\mathcal{B}(\alpha_{vk}^{i*}, \alpha_{vk}^{i*+})}{\mathcal{B}(\alpha_{vk}, \alpha_{vk}^+)}} & y_{vk}^i = 0, \end{cases}$$

$$S_{vk}^{i*} = \psi(\alpha_{vk}^{i*+}) - \psi(\alpha_{vk}^{i*} + \alpha_{vk}^{i*+})$$

$$R_{vk}^{i*} = \psi(\alpha_{vk}^{i*}) - \psi(\alpha_{vk}^{i*+})$$

- $\alpha_{vk}^{i*} = \alpha_{vk} + y_{vk}^i, \quad \alpha_{vk}^{i*+} = \sum_{j=k+1}^{K_v} \alpha_{vk}^{i*}, \quad \psi(\cdot)$ is the digamma function.

E-step: Q-function

- We can rewrite $Q(\theta)$ as:

$$Q(\theta) = \sum_{i=1}^n \sum_{v \in \mathcal{V}} \sum_{k=1}^{K_v-1} \{Q_1(\pi_{vk}, \delta_{vk}^i) + Q_2(\alpha_{vk}, \alpha_{vk}^+, R_{vk}^{i*}, S_{vk}^{i*})\}$$

- where

$$Q_1(\pi_{vk}, \delta_{vk}^i) = \delta_{vk}^{i*} \log(\pi_{vk}) + (1 - \delta_{vk}^{i*}) \log(1 - \pi_{vk})$$

and

$$Q_2(\alpha_{vk}, \alpha_{vk}^+, R_{vk}^{i*}, S_{vk}^{i*}) = (1 - \delta_{vk}^{i*}) \{-\log \mathcal{B}(\alpha_{vk}, \alpha_{vk}^+) \\ + (\alpha_{vk} - 1) R_{vk}^{i*} + (\alpha_{vk}^+ - 1) S_{vk}^{i*}\}$$

- Maximize $Q(\theta)$ with respect to θ
- The parameters in Q_1 and Q_2 can be optimized separately
- Model depends on ordering of OTUs at each internal node (matching problem)
- Fit separate ZIDM model for each possible ordering of $|C_v|$ taxa, and select the best fitted model.
- Computational cost is $O(\sum_{v \in \mathcal{V}} |C_v|!)$ compared to $O(|\mathcal{L}|!)$ for GDM or ZIDM models

Binary trees

- From now on, consider binary trees
- i.e. $|C_v| = 2$ for all $v \in \mathcal{V}$
- Dirichlet, multinomial, dirichlet multinomial, and zero inflated dirichlet multinomial are reduced to beta, binomial, beta-binomial and zero-inflated beta-binomial.
- The ZIDTM is then a product of ZIBBs

Posterior mean transformation

- Estimate underlying proportions from Bayesian perspective (since Dirichlet is conjugate to multinomial)
- Consider the posterior mean for the Dirichlet Multinomial:

$$E_{DM}(p_k|\mathbf{y}) = \frac{\alpha_k + y_k}{\sum_{j=1}^K (y_j + \alpha_j)}$$

- We estimate unknown parameters by maximizing data likelihood
- Use the estimated parameters as "pseudo data" and combine with real observed data.
- This produces non-zero proportions for zero counts
- Alternative to using a pseudo-count for sample proportions

Solving posterior mean transformation

- When we have a binary tree, there is an explicit closed-form solution to the zero-inflation
- At each internal node $v \in \mathcal{V}$,

$$E_{BB}(p_{v1}|\mathbf{y}_v) = \frac{\alpha_{v1} + y_{v1}}{\alpha_{v1} + \alpha_{v2} + y_{v1} + y_{v2}},$$

and

$$E_{ZIBB}(p_{v1}|\mathbf{y}_v) = \frac{(1 - B_{v0})B_{v1}}{B_{v0}\mathcal{B}(\alpha_{v1}, \alpha_{v2}) + (1 - B_{v0})B_{v2}},$$


- $B_{v0} = \pi I(y_{v1} = 0)$
- $B_1 = \mathcal{B}(1 + \alpha_{v1} + y_{v1}, \alpha_{v2} + y_{v2})$
- $B_2 = \mathcal{B}(\alpha_{v1} + y_{v1}, \alpha_{v2} + y_{v2})$

Likelihood-ratio test & count transformation

- Having a correct model specification at internal nodes effects quality of posterior estimates.
- Perform a two-stage likelihood-ratio test
- First, assume count data at $v \in \mathcal{V}$ are not zero-inflated, fit beta-binomial model and test for over-dispersion.
- Counts at nodes without over-dispersion are transformed into sample proportions after adding constant of .5, equivalent to using E_{BB} with $\alpha_{v1} = \alpha_{v2} = 0.5$
- Second, Nodes with over-dispersion are refit using a ZIBB model, and tested again for zero-inflation
- Counts are then transformed using $E_{ZIBB}(p_{v1}|\mathbf{y}_v)$.
- If there is no zero-inflation, counts are transformed using $E_{BB}(p_{v1}, \mathbf{y}_v)$
- $\hat{\mathbf{p}}_v = (\hat{p}_{v1}, 1 - \hat{p}_{v1})^T$ is the posterior estimate

Path-level definitions - Phylogeny-aware normalization

- To ensure that the normalization is 'phylogeny-aware', we consider the path level instead of focussing on individual internal nodes.
- Define \mathcal{A}_u as the ancestor node set of u for each node $u \in \mathcal{L} \cup \mathcal{V}$ that contains all internal nodes in the path from the root node to u
- \mathcal{L}_u is the set of leaves in the same path



- $\mathcal{A}_1 = \mathcal{A}_2 = \{11, 12, 13\}$,
 $\mathcal{A}_3 = \{11, 12\}$, $\mathcal{L}_{11} = \mathcal{L}$, $\mathcal{L}_{12} = \{1, 2, 3\}$, $\mathcal{L}_{13} = \{1, 2\}$
- Define $q_u = \prod_{v \in \mathcal{A}_u} \hat{p}_v$
- \mathcal{U} is the set of nodes such that $\cup_{u \in \mathcal{U}} \mathcal{L}_u = \mathcal{L}$, $\mathcal{L}_u \cap \mathcal{L}_{u'} = \emptyset$ when $u \neq u'$.
- Then, $\sum_{u \in \mathcal{U}} q_u = 1$ and $\sum_{l \in \mathcal{L}} q_l = 1$

- Detect differentially abundant OTUs at ecosystem level
- Similar to ANCOM
- Difference is how we use the phylogenetic tree to incorporate this information
- Tests differences between two groups
- Speed up ANCOM by using phylogenetic tree to construct adaptive log-ratios

Hypothesis tests - ANCOM

- Test the hypothesis

$$H_{0l} : \log \mu_l^1 = \log \mu_l^2$$

for $l \in \mathcal{L} = \{1, \dots, K\}$, and where μ_l^g is the mean absolute abundance of the l th OTU from g th group, and $g = 1, 2$

- In ANCOM, we test $K - 1$ hypothesis using each OTU as the denominator in the additive log ratio

$$H_{0lm} : \log(\mu_l^1 / \mu_m^1) = \log(\mu_l^2 / \mu_m^2)$$

for all $m \neq l$

- ANCOM counts the number of rejections across the $K - 1$ hypothesis tests, and compares to the empirical distribution of counts to determine if the l th OTU is differentially abundant.
- If K is high, ANCOM suffers from a high FDR. Additionally the computational burden is high

adaANCOM log-ratios & hypothesis tests

- Assume that abundance difference on the log scale at an internal node is passed down to child nodes
- adaANCOM is comprised of a two-step process
 - Test the internal node-level hypothesis

$$H_{0v} : \log(\mu_{v1}^1 / \mu_{v2}^1) = \log(\mu_{v1}^2 / \mu_{v2}^2)$$

for $v \in \mathcal{V}$, with μ_{v1}^g and μ_{v2}^g as mean absolute abundances of children of v from the g th group.

For a given α , $\mathcal{D}_{\mathcal{V}}$ is the set of internal nodes which the hypotheses are rejected

- For each leaf node $l \in \mathcal{L}$, calculate the log-ratio
Define ref_l as the sibling node of l if $\mathcal{A}_l \cap \mathcal{D}_{\mathcal{V}} = \emptyset$ and the child node of v not in \mathcal{A}_l otherwise, when v is the node in $\mathcal{A}_l \cap \mathcal{D}_{\mathcal{V}}$ closest to the root node.


The null hypothesis is:

$$H_{0l}^{ada} : \log(\mu_l^1 / \mu_{ref_l}^1) = \log(\mu_l^2 / \mu_{ref_l}^2)$$

- adaANCOM rejects the null hypothesis H_{0l} if H_{0l}^{ada} is rejected

adaANCOM example

- Assume $\mathcal{D}_Y = 17$.
- $ref_1 = 2, ref_2 = 1$
 $ref_3 = 13, ref_4 = 16,$
 $ref_5 = 6, ref_6 = 5,$
 $ref_7 = ref_8 = 19,$
 $ref_9 = ref_{10} = 18$



Accounting for log-ratios of zero y_u counts

- H_{0v} and H_{0l}^{ada} are based on the log-ratios of q_u s for $u \in \mathcal{L} \cup \mathcal{V}$, but when the original y_u counts have many zero counts, log-ratios give abnormal results, and test statistics are sensitive to these abnormal values
- For internal node v , define ϕ_v as the maximum of

$$|\log(\hat{p}_{v1}/\hat{p}_{v2})|$$

that have $y_{v1} > 0$ and $y_{v2} > 0$.

- Data with $y_{v1} = 0$ or $y_{v2} = 2$ are removed if $|\log(\hat{p}_{v1}/\hat{p}_{v2})| > \phi_v$

Comparisons to ANCOM

- Similar to ANCOM, adaANCOM uses relative abundance data, constructs log-ratios, and performs t-tests or Wilcoxon-rank-sum tests.
- The advantage of adaANCOM is the computation time, as the number of tests is reduced by a factor of $|\mathcal{L}|$.
- adaANCOM also has an advantage of interpretability since we are guided by the tree, so we result in both DA leaves and nodes.
- Additionally, adaANCOM controls FDR better

Algorithm 1: adaANCOM

Input: A binary tree $\mathcal{T} = (\ell, \mathcal{V})$, posterior-mean-transformed data $\{\hat{p}_v, v \in \mathcal{V}\}$, group information, and a testing procedure;

Output: DA internal nodes $\mathcal{D}_\mathcal{V}$, and DA leaf nodes \mathcal{D}_ℓ ;

Step 1:

Set $\mathcal{D}_\mathcal{V} = \emptyset$, **for** $v \in \mathcal{V}$ **do**

 Construct the log-ratios $\log(\hat{p}_{v1}/\hat{p}_{v2})$, and remove the outliers;

 Test H_{0v} , and if rejected, update $\mathcal{D}_\mathcal{V} = \mathcal{D}_\mathcal{V} \cup \{v\}$;

end

Step 2:

Set $\mathcal{D}_\ell = \emptyset$, **for** $l \in \ell$ **do**

 Search ref_l , compute q_l and q_{ref_l} ;

 Construct the log-ratios $\log(q_l/q_{ref_l})$, and remove the outliers;

 Test H_{0l}^{ada} , and if rejected, update $\mathcal{D}_\ell = \mathcal{D}_\ell \cup \{l\}$;

end